

Lectures

Stat 154/254: Statistical Machine Learning

Lectures marked with ★ are still in progress.



Introduction

- Class introduction

Unit 1: Regression in a machine learning setting

- ★ Linear regression in matrix notation (review)
- ★ Linear regression in an ML setting
- ★ Basis expansions and splines for linear regression
- ★ L1 and L2 penalization

Unit 2: Classification

- ★ Empirical risk minimization for classification
- ★ Trading off false positives and negatives
- ★ Perceptron algorithm (**bonus content**)

Unit 3: Risk and complexity

- ★ Asymptotics and bias–variance tradeoff in linear regression
- ★ Regret decomposition
- Cross validation
 - ★ Held-out data and cross-validation
 - ★ Simulated bias–variance tradeoff example
 - ★ Cross-validation cautionary note
- Generalization bounds
 - ★ Generalization error and uniform laws
 - ★ Uniform laws for smooth function classes
 - ★ VC dimension and zero–one loss

Unit 4: Trees and weak learners

- ★ Classification and regression trees
- ★ Bagging weak learners
- ★ Boosting weak learners

Unit 5: Neural networks and optimization

- [★ Optimization notes](#)

Prerequisites

Mathematical statistics, probability, linear algebra, multivariate calculus, and Python are all prerequisites.

For assistance with mathematics prerequisites, the [Student Learning Center](#) provides additional assistance.

A summary of the required linear algebra skills can be found [here](#).